

```
# Load the Drive module:
from google.colab import drive
```

```
drive.mount("/content/drive")
! ls "/content/drive/MyDrive"
```

```
Mounted at /content/drive
'Build AI for a Better Society'      'Learning Insights (Bachelor)'
'Build Your Dream Team '            'Mission Identification (Bachelor)'
```

'Calibration Phase - Onboarding Time! '	'Orientation Group Challenge - 2023'
cars_dataset.csv	pairplot.pdf
'Colab Notebooks'	pairplot.png
correlation_heatmap.pdf	'Strategy & Global Markets'
'Foundations for Tech Impact'	'Sustainability Foundations (Bx)'
'Introduction to Coding and AI (BX)'	'Technology Revolutions'

```
import os
print(os.getcwd())
print(os.listdir())
print(os.listdir("drive/MyDrive"))
```

```
/content
['.config', 'drive', 'sample_data']
['Learning Insights (Bachelor)', 'Mission Identification (Bachelor)', 'Build Your Dream Team ', 'Introduction to Coding and AI (BX)']
```

```
#Milestone 3 Code
import pandas as pd
```

```
df = pd.read_csv("/content/drive/MyDrive/cars_dataset.csv")
df.head()
```

```
0
```

	uuid	Make	Model	Vehicle Class	Engine Size(L)	Cylinders	Transmis
0	eb24f214-f18b-11ec-a33e-acde48001122	CHEVROLET	Suburban 4WD	SUV - STANDARD	5.3	8	
1	eb2525fe-f18b-11ec-a33e-acde48001122	CHEVROLET	CAMARO	SUBCOMPACT	3.6	6	
2	eb252d2e-f18b-11ec-a33e-acde48001122	LEXUS	GS 350	MID-SIZE	3.5	6	

```
df = pd.read_csv("/content/drive/MyDrive/cars_dataset.csv")
```

```
#Classification of datatypes
df.dtypes
```

```
uuid          object
Make          object
Model         object
Vehicle Class object
Engine Size(L) float64
Cylinders     int64
Transmission  object
Fuel Type     object
Fuel Consumption City (L/100 km) float64
Fuel Consumption Hwy (L/100 km) float64
Fuel Consumption Comb (L/100 km) float64
Fuel Consumption Comb (mpg)      int64
CO2 Emissions(g/km)             int64
dtype: object
```

```
df = pd.read_csv("/content/drive/MyDrive/cars_dataset.csv")
```

```
# Count the number of columns
num_columns = df.shape[1]
```

```
# Print the result
print("Number of columns:", num_columns)
```

```
Number of columns: 13
```

```
df = pd.read_csv("/content/drive/MyDrive/cars_dataset.csv")
```

```
# Get the dimensions (number of rows and columns)
dimensions = df.shape
```

```
# Print the result
print("Dimensions:", dimensions)
```

```
↳ Dimensions: (7385, 13)
```

```
df = pd.read_csv("/content/drive/MyDrive/cars_dataset.csv")
```

```
df["Make"]
```

```
↳
0      CHEVROLET
1      CHEVROLET
2      LEXUS
3      LEXUS
4      PORSCHE
...
7380     HONDA
7381  MERCEDES-BENZ
7382     HYUNDAI
7383     PORSCHE
7384      FORD
Name: Make, Length: 7385, dtype: object
```

```
df = pd.read_csv("/content/drive/MyDrive/cars_dataset.csv")
```

```
df['Make'].unique()
```

```
↳ array(['CHEVROLET', 'LEXUS', 'PORSCHE', 'BMW', 'AUDI', 'NISSAN',
        'LINCOLN', 'VOLKSWAGEN', 'TOYOTA', 'ROLLS-ROYCE', 'MERCEDES-BENZ',
        'HONDA', 'HYUNDAI', 'KIA', 'GMC', 'ACURA', 'FORD', 'ASTON MARTIN',
        'JEEP', 'JAGUAR', 'DODGE', 'MAZDA', 'VOLVO', 'SUBARU', 'RAM',
        'CADILLAC', 'LAMBORGHINI', 'MINI', 'MITSUBISHI', 'MASERATI',
        'SCION', 'CHRYSLER', 'LAND ROVER', 'BENTLEY', 'INFINITI',
        'GENESIS', 'FIAT', 'BUICK', 'BUGATTI', 'ALFA ROMEO', 'SRT',
        'SMART'], dtype=object)
```

```
df = pd.read_csv("/content/drive/MyDrive/cars_dataset.csv")
```

```
df['Make'].value_counts()
```

```
↳
FORD      628
CHEVROLET 588
BMW       527
MERCEDES-BENZ 419
PORSCHE   376
TOYOTA    330
GMC       328
AUDI      286
NISSAN    259
JEEP      251
DODGE     246
KIA       231
HONDA     214
HYUNDAI   210
MINI      204
VOLKSWAGEN 197
MAZDA     180
LEXUS     178
JAGUAR    160
CADILLAC  158
SUBARU    140
VOLVO     124
INFINITI  108
BUICK     103
RAM       97
LINCOLN   96
MITSUBISHI 95
CHRYSLER  88
LAND ROVER 85
FIAT      73
ACURA    72
MASERATI  61
ROLLS-ROYCE 50
ASTON MARTIN 47
BENTLEY   46
LAMBORGHINI 41
ALFA ROMEO 30
GENESIS   25
SCION     22
SMART     7
BUGATTI   3
```

```
SRT                2
Name: Make, dtype: int64

df = pd.read_csv("/content/drive/MyDrive/cars_dataset.csv")

df.info()
<> <class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   uuid                                7385 non-null   object
1   Make                               7385 non-null   object
2   Model                              7385 non-null   object
3   Vehicle Class                      7385 non-null   object
4   Engine Size(L)                     7385 non-null   float64
5   Cylinders                           7385 non-null   int64
6   Transmission                        7385 non-null   object
7   Fuel Type                           7385 non-null   object
8   Fuel Consumption City (L/100 km)    7385 non-null   float64
9   Fuel Consumption Hwy (L/100 km)     7385 non-null   float64
10  Fuel Consumption Comb (L/100 km)    7385 non-null   float64
11  Fuel Consumption Comb (mpg)         7385 non-null   int64
12  CO2 Emissions(g/km)                7385 non-null   int64
dtypes: float64(4), int64(3), object(6)
memory usage: 750.2+ KB

df = pd.read_csv("/content/drive/MyDrive/cars_dataset.csv")

df.corr(numeric_only=True)
<>
           Engine Size(L)  Cylinders  Fuel Consumption City (L/100 km)  Fuel Consumption Hwy (L/100 km)  Fuel Consumption Comb (L/100 km)  Fuel Consumption Comb (mpg)
Engine Size(L)      1.000000      0.927653      0.831379      0.761526      0.817060      -0.757
Cylinders           0.927653      1.000000      0.800702      0.715252      0.780534      -0.719
Fuel Consumption City (L/100 km)  0.831379      0.800702      1.000000      0.948180      0.993810      -0.927
Fuel Consumption Hwy (L/100 km)  0.761526      0.715252      0.948180      1.000000      0.977000      -0.800
Fuel Consumption Comb (L/100 km) 0.817060      0.780534      0.993810      0.977000      1.000000      -0.600
Fuel Consumption Comb (mpg)     -0.757      -0.719      -0.927      -0.800      -0.600      1.000

df = pd.read_csv("/content/drive/MyDrive/cars_dataset.csv")

df.describe()
<>
           Engine Size(L)  Cylinders  Fuel Consumption City (L/100 km)  Fuel Consumption Hwy (L/100 km)  Fuel Consumption Comb (L/100 km)  Fuel Consumption Comb (mpg)  count
count  7385.000000      7385.000000      7385.000000      7385.000000      7385.000000      7385.000000      7385.000000
mean    3.160068        5.615030      12.556534      9.041706      10.975071      27.481652
std     1.354170        1.828307      3.500274      2.224456      2.892506      7.231879
min     0.900000        3.000000      4.200000      4.000000      4.100000      11.000000
25%     2.000000        4.000000      10.100000      7.500000      8.900000      22.000000
50%     3.000000        6.000000      12.100000      8.700000      10.600000      27.000000

df = pd.read_csv("/content/drive/MyDrive/cars_dataset.csv")

# Compute the distribution for all columns
distributions = {}
for column in df.columns:
    distribution = df[column].value_counts()
    distributions[column] = distribution

# Print the distributions
for column, distribution in distributions.items():
    print("Distribution of", column)
    print(distribution)
    print()
<> Distribution of uuid
eb24f214-f18b-11ec-a33e-acde48001122    1
```

```

eb733a46-f18b-11ec-a33e-acde48001122    1
eb73335c-f18b-11ec-a33e-acde48001122    1
eb732fec-f18b-11ec-a33e-acde48001122    1
eb732c72-f18b-11ec-a33e-acde48001122    1
..
eb4e8e3a-f18b-11ec-a33e-acde48001122    1
eb4e8a8e-f18b-11ec-a33e-acde48001122    1
eb4e8638-f18b-11ec-a33e-acde48001122    1
eb4e8228-f18b-11ec-a33e-acde48001122    1
eb98f772-f18b-11ec-a33e-acde48001122    1
Name: uuid, Length: 7385, dtype: int64

```

```

Distribution of Make
FORD      628
CHEVROLET 588
BMW       527
MERCEDES-BENZ 419
PORSCHE   376
TOYOTA    330
GMC       328
AUDI      286
NISSAN    259
JEEP      251
DODGE     246
KIA       231
HONDA     214
HYUNDAI   210
MINI      204
VOLKSWAGEN 197
MAZDA     180
LEXUS     178
JAGUAR    160
CADILLAC  158
SUBARU    140
VOLVO     124
INFINITI  108
BUICK     103
RAM       97
LINCOLN   96
MITSUBISHI 95
CHRYSLER  88
LAND ROVER 85
FIAT      73
ACURA    72
MASERATI  61
ROLLS-ROYCE 50
ASTON MARTIN 47
BENTLEY   46
LAMBORGHINI 41
ALFA ROMEO 30
GENESIS   25
SCION     22
SMART     7
BUGATTI   3
SRT       2
Name: Make, dtype: int64

```

```

!pip install pandas
!pip install seaborn

```

```

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (1.5.3)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2022.7.1)
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.10/dist-packages (from pandas) (1.22.4)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1->pandas) (1.16.0)
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (0.12.2)
Requirement already satisfied: numpy!=1.24.0,>=1.17 in /usr/local/lib/python3.10/dist-packages (from seaborn) (1.22.4)
Requirement already satisfied: pandas>=0.25 in /usr/local/lib/python3.10/dist-packages (from seaborn) (1.5.3)
Requirement already satisfied: matplotlib!=3.6.1,>=3.1 in /usr/local/lib/python3.10/dist-packages (from seaborn) (3.7.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.0.7)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (4.22.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.4.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (23.1)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (8.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (3.1.0)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.25->seaborn) (2022.7.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.1) (1.16.0)

```

```

import seaborn as sns
import matplotlib.pyplot as plt

# Read the CSV file into a pandas DataFrame
file_path = '/content/drive/MyDrive/cars_dataset.csv'
df = pd.read_csv(file_path)


# Calculate the correlation matrix
corr_matrix = df.corr()

# Create a heatmap using seaborn
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Heatmap')

# Save the heatmap as a PDF file
plt.savefig('/content/drive/MyDrive/correlation_heatmap.pdf', format='pdf')

```

<ipython-input-16-e05a3341e700>:9: FutureWarning: The default value of numeric_only is deprecated



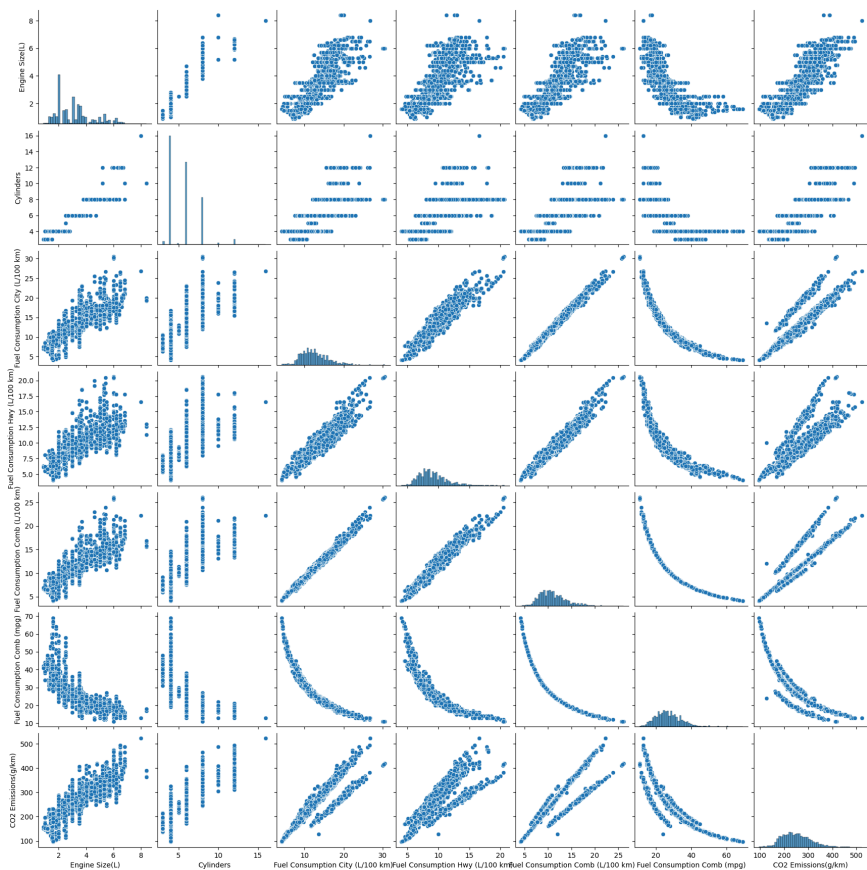
```

import pandas as pd
import seaborn as sns

sns.pairplot(df)

# Save the pair plot as a PNG file
plt.savefig('/content/drive/MyDrive/pairplot.png', dpi=300)

```



#Milestone 4 Code

```
import os
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

df = pd.read_csv("/content/drive/MyDrive/cars_dataset.csv")

# -----
# Import the data
# -----
df = pd.read_csv('/content/drive/MyDrive/cars_dataset.csv')
print(df.info())

# Get Cylinders count
dMean_cylinders = df['Cylinders'].mean()
print('// complete ..... data model cylinders mean: ', dMean_cylinders)

# Generate a small image for slides ;- )
fig, ax = plt.subplots()
ax.spines['right'].set_visible(False)
ax.spines['top'].set_visible(False)
ax.yaxis.set_ticks_position('left')
ax.xaxis.set_ticks_position('bottom')
sns.scatterplot(x='Fuel Consumption Comb (L/100 km)',
                y='CO2 Emissions(g/km)',
                data=df, alpha=0.5, color='grey')# palette='Grays')
plt.savefig('fig_scatter_Fuel_CO2.pdf')
plt.close()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   uuid                                7385 non-null   object
1   Make                                7385 non-null   object
2   Model                               7385 non-null   object
3   Vehicle Class                       7385 non-null   object
4   Engine Size(L)                       7385 non-null   float64
5   Cylinders                           7385 non-null   int64
6   Transmission                         7385 non-null   object
7   Fuel Type                           7385 non-null   object
8   Fuel Consumption City (L/100 km)     7385 non-null   float64
9   Fuel Consumption Hwy (L/100 km)      7385 non-null   float64
10  Fuel Consumption Comb (L/100 km)     7385 non-null   float64
11  Fuel Consumption Comb (mpg)          7385 non-null   int64
12  CO2 Emissions(g/km)                 7385 non-null   int64
dtypes: float64(4), int64(3), object(6)
memory usage: 750.2+ KB
None
// complete ..... data model cylinders mean:  5.615030467163169
```

```
# -----
# Import the data
# -----
df = pd.read_csv('/content/drive/MyDrive/cars_dataset.csv')
print(df.info())

# Get Cylinders count
dMean_cylinders = df['Cylinders'].mean()
print('// complete ..... data model cylinders mean: ', dMean_cylinders)

# Generate a small image for slides ;- )
fig, ax = plt.subplots()
ax.spines['right'].set_visible(False)
ax.spines['top'].set_visible(False)
ax.yaxis.set_ticks_position('left')
ax.xaxis.set_ticks_position('bottom')
sns.scatterplot(x='CO2 Emissions(g/km)',
                y='Cylinders',
                data=df, alpha=0.5, color='grey')# palette='Grays')
plt.savefig('fig_scatter_Cylinders_Fuel_CO2.pdf')
plt.close()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 13 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   uuid                                     7385 non-null   object
1   Make                                    7385 non-null   object
2   Model                                  7385 non-null   object
3   Vehicle Class                          7385 non-null   object
4   Engine Size(L)                         7385 non-null   float64
5   Cylinders                              7385 non-null   int64
6   Transmission                           7385 non-null   object
7   Fuel Type                              7385 non-null   object
8   Fuel Consumption City (L/100 km)       7385 non-null   float64
9   Fuel Consumption Hwy (L/100 km)        7385 non-null   float64
10  Fuel Consumption Comb (L/100 km)       7385 non-null   float64
11  Fuel Consumption Comb (mpg)            7385 non-null   int64
12  CO2 Emissions(g/km)                   7385 non-null   int64
dtypes: float64(4), int64(3), object(6)
memory usage: 750.2+ KB
None
// complete ..... data model cylinders mean:  5.615030467163169
```

```
# Selecting the desired columns
columns_of_interest = ['Cylinders', 'Fuel Consumption Comb (L/100 km)', 'CO2 Emissions(g/km)']
subset_df = df[columns_of_interest]
```

```
# Applying the describe() method on the subset DataFrame
description = subset_df.describe()
```

```
# Printing the description
print(description)
```

```
<class 'pandas.core.frame.DataFrame'>
Cylinders  Fuel Consumption Comb (L/100 km)  CO2 Emissions(g/km)
count      7385.000000                      7385.000000          7385.000000
mean         5.615030                      10.975071          250.584699
std          1.828307                      2.892506           58.512679
min           3.000000                      4.100000           96.000000
25%           4.000000                      8.900000          208.000000
50%           6.000000                     10.600000          246.000000
75%           6.000000                     12.600000          288.000000
max          16.000000                     26.100000          522.000000
```



```

from sklearn.linear_model import LinearRegression

# Import the data
df = pd.read_csv('/content/drive/MyDrive/cars_dataset.csv')
print(df.info())

# Get Cylinders count
dMean_cylinders = df['Cylinders'].mean()
print('// complete ..... data model cylinders mean: ', dMean_cylinders)

# Generate a scatter plot
fig, ax = plt.subplots()
ax.spines['right'].set_visible(False)
ax.spines['top'].set_visible(False)
ax.yaxis.set_ticks_position('left')
ax.xaxis.set_ticks_position('bottom')
sns.scatterplot(x='CO2 Emissions(g/km)', y='Cylinders', data=df, alpha=0.5, color='grey')

# Perform linear regression
regression_model = LinearRegression()
X = df[['CO2 Emissions(g/km)']] # Input feature
y = df['Cylinders'] # Target variable
regression_model.fit(X, y)

# Plot the linear regression line
plt.plot(X, regression_model.predict(X), color='red')

plt.savefig('fig_scatter_Linear_Regression_2_Cylinders_Fuel_CO2.pdf')
plt.close()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 13 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   uuid                                7385 non-null   object
 1   Make                                7385 non-null   object
 2   Model                               7385 non-null   object
 3   Vehicle Class                       7385 non-null   object
 4   Engine Size(L)                      7385 non-null   float64
 5   Cylinders                           7385 non-null   int64
 6   Transmission                        7385 non-null   object
 7   Fuel Type                           7385 non-null   object
 8   Fuel Consumption City (L/100 km)    7385 non-null   float64
 9   Fuel Consumption Hwy (L/100 km)     7385 non-null   float64
10   Fuel Consumption Comb (L/100 km)    7385 non-null   float64
11   Fuel Consumption Comb (mpg)         7385 non-null   int64
12   CO2 Emissions(g/km)                 7385 non-null   int64
dtypes: float64(4), int64(3), object(6)
memory usage: 750.2+ KB
None
// complete ..... data model cylinders mean:  5.615030467163169

```

```

# Import the data
df = pd.read_csv('/content/drive/MyDrive/cars_dataset.csv')
print(df.info())

# Get Cylinders count
dMean_cylinders = df['Cylinders'].mean()
print('// complete ..... data model cylinders mean: ', dMean_cylinders)

# Generate a scatter plot
fig, ax = plt.subplots()
ax.spines['right'].set_visible(False)
ax.spines['top'].set_visible(False)
ax.yaxis.set_ticks_position('left')
ax.xaxis.set_ticks_position('bottom')
sns.scatterplot(x='Fuel Consumption Comb (L/100 km)', y='CO2 Emissions(g/km)', data=df, alpha=0.5, color='grey')

# Perform linear regression
regression_model = LinearRegression()
X = df[['Fuel Consumption Comb (L/100 km)']] # Input feature
y = df['CO2 Emissions(g/km)'] # Target variable
regression_model.fit(X, y)

# Plot the linear regression line
plt.plot(X, regression_model.predict(X), color='red')

plt.savefig('fig_scatter_Linear_Regression_2CO2_Emissions_Fuel_CO2.pdf')
plt.close()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384

```

```
Data columns (total 13 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   uuid                                     7385 non-null   object
1   Make                                    7385 non-null   object
2   Model                                   7385 non-null   object
3   Vehicle Class                           7385 non-null   object
4   Engine Size(L)                           7385 non-null   float64
5   Cylinders                                7385 non-null   int64
6   Transmission                             7385 non-null   object
7   Fuel Type                                7385 non-null   object
8   Fuel Consumption City (L/100 km)         7385 non-null   float64
9   Fuel Consumption Hwy (L/100 km)          7385 non-null   float64
10  Fuel Consumption Comb (L/100 km)          7385 non-null   float64
11  Fuel Consumption Comb (mpg)               7385 non-null   int64
12  CO2 Emissions(g/km)                       7385 non-null   int64
dtypes: float64(4), int64(3), object(6)
memory usage: 750.2+ KB
None
// complete ..... data model cylinders mean: 5.615030467163169
```

```
#Residual Plot, r-squared and RMSE for Evaluation Metrics (Capstone)
```

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
```

```
df = pd.read_csv('/content/drive/MyDrive/cars_dataset.csv')
```

```
# Load data
```

```
X = df[['Cylinders']].values
```

```
y = df['CO2 Emissions(g/km)'].values
```

```
# Split data into train and test sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
# Train linear regression model
```

```
model = LinearRegression()
```

```
model.fit(X_train, y_train)
```

```
# Model evaluation metrics
```

```
r_squared = model.score(X_test, y_test)
```

```
print(f"R-Squared: {r_squared}")
```

```
y_pred = model.predict(X_test)
```

```
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
```

```
print(f"RMSE: {rmse}")
```

```
# Residual plot
```

```
residuals = y_test - y_pred
```

```
plt.scatter(y_pred, residuals)
```

```
plt.xlabel("Predicted CO2")
```

```
plt.ylabel("Residuals")
```

```
# Feature importance
```

```
print("Slope", model.coef_)
```


```
print("Intercept", model.intercept_)
```

```
→ R-Squared: 0.7061506959702248
```

```
RMSE: 31.655771694980565
```

```
Slope [26.57554309]
```

```
Intercept 101.26577344379822
```



```

#Milestone 6 Code
# -----
# Pre-Processing
# -----
# Data pre-processing
# Let's check the dataset for missing values.
# @code: 0, or 'index': Drop rows which contain missing values.
# Let's see where the Null values are.
# Let's see the data shape and NaN values.
# This will give number of NaN values in every column.
df_null_values = df.isnull().sum()
print('NANs?', df_null_values)

# Show missing values in a figure
# plt.figure(figsize=(15,5))
# sns.heatmap(df.isnull(), cbar=False, yticklabels=False, cmap='Greys')
# plt.xticks(rotation=45, fontsize=6)
# plt.tight_layout()
# plt.savefig('fig_MissingValues.pdf')
# plt.close()

# Drop all rows with NaN.
df = df.dropna(axis=0)
df_null_values = df.isnull().sum()
print('NANs_After_Update?', df_null_values)
print('// complete ..... Pre-Processing')

```

```

→ NANs? uuid          0
  Make                0
  Model               0
  Vehicle Class       0
  Engine Size(L)      0
  Cylinders           0
  Transmission        0
  Fuel Type           0
  Fuel Consumption City (L/100 km)  0
  Fuel Consumption Hwy (L/100 km)  0
  Fuel Consumption Comb (L/100 km)  0
  Fuel Consumption Comb (mpg)       0
  CO2 Emissions(g/km)              0
  dtype: int64
  NANs_After_Update? uuid          0
    Make                0
    Model               0
    Vehicle Class       0
    Engine Size(L)      0
    Cylinders           0
    Transmission        0
    Fuel Type           0
    Fuel Consumption City (L/100 km)  0
    Fuel Consumption Hwy (L/100 km)  0
    Fuel Consumption Comb (L/100 km)  0
    Fuel Consumption Comb (mpg)       0
    CO2 Emissions(g/km)              0
    dtype: int64
  // complete ..... Pre-Processing

```

```

# Check for missing values
missing_values = df.isnull().sum()
print("Missing values:\n", missing_values)

```

```

# Check datatypes
datatypes = df.dtypes
print("Datatypes:\n", datatypes)

```

```

→ Missing values:
  uuid          0
  Make          0

```