

```
# Load the Drive module:
from google.colab import drive

# Authenticate user manually
from google.colab import auth
auth.authenticate_user()

# Mount the drive
drive.mount('/content/drive')

# List the contents of MyDrive
!ls "/content/drive/MyDrive"
```



Mounted at /content/drive

'Additional Factors Summary Statistics.gsheet'	'Introduction to Coding and AI (BX)'
'Build AI for a Better Society'	'Learning Insights (Bachelor)'
'Build Your Dream Team '	'Mission Identification (Bachelor)'
'Business and Society'	'Orientation Group Challenge - 2023'
'Calibration Phase - Onboarding Time! '	pairplot_df-encoded.png
cars_dataset.csv	pairplot_machinelearning.png
'Colab Notebooks'	pairplot.pdf
correlation_heatmap.pdf	pairplot.png
'Data Products'	'Problem Solving with Data'
'Data Storytelling for Impact'	'Strategy & Global Markets'
final_dataset.csv	'Sustainability Foundations (Bx)'
'Foundations for Tech Impact'	'Technology Revolutions'

✓ Machine Learning and Deep Learning Notebook

- **Task:** Predict the CO2 emissions of cars utilizing a neural network and a supervised machine learning model.
- **Dataset:** final_dataset.csv

```
#Import Pandas
import pandas as pd

# Importing data
# Creating DataFrame
# Storing DataFrame in a variable
df = pd.read_csv("/content/drive/MyDrive/final_dataset.csv")
```

✓ Columns of the Dataset

```
# Get the column names
columns = df.columns

# Create a Markdown table string
markdown_table = "| Column Index | Column Name |\n|-----|-----|\n"
for i, col in enumerate(columns):
    markdown_table += f"| {i} | {col} |\n"

# Display the table in a text cell using Markdown
from IPython.display import display, Markdown

display(Markdown(markdown_table))
```



Column Index	Column Name
0	uuid
1	Make
2	Model
3	Vehicle Class
4	Engine Size(L)
5	Cylinders
6	Transmission
7	Fuel Type
8	Fuel Consumption City (L/100 km)
9	Fuel Consumption Hwy (L/100 km)
10	Fuel Consumption Comb (L/100 km)
11	Fuel Consumption Comb (mpg)
12	CO2 Emissions(g/km)

```
# Dropping unnecessary columns
columns_to_drop = ['uuid', 'Make', 'Vehicle Class', 'Transmission', 'Fuel Type', 'Model']
df.drop(columns=columns_to_drop, inplace=True)

# Importing train_test_split function from scikit-learn
from sklearn.model_selection import train_test_split

# Define X and y
X = df.drop(['CO2 Emissions(g/km)'], axis=1) # Drop the target column 'CO2 Emissions(g/km)'
y = df['CO2 Emissions(g/km)'] # Target column 'CO2 Emissions(g/km)' as y

# Splitting the data into training and testing sets with a test size of 20% and random state
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

```
# Print the first few rows of the training set
print("Training set:")
print(X_train.head())
print(y_train.head())

# Print the first few rows of the testing set
print("\nTesting set:")
print(X_test.head())
print(y_test.head())
```



Training set:

	Engine Size(L)	Cylinders	Fuel Consumption City (L/100 km)	\
1095	3.5	6	12.7	
3538	5.7	8	16.1	
2377	3.0	6	11.8	
3414	1.4	4	8.7	
880	3.6	6	12.3	

	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)	\
1095	9.5	11.3	
3538	11.0	13.8	
2377	8.4	10.3	
3414	7.1	8.0	
880	8.2	10.4	

	Fuel Consumption Comb (mpg)
1095	25
3538	20
2377	27
3414	35
880	27

1095	264
3538	325
2377	237
3414	186
880	246

Name: CO2 Emissions(g/km), dtype: int64

Testing set:

	Engine Size(L)	Cylinders	Fuel Consumption City (L/100 km)	\
6307	2.0	4	10.2	
5036	1.5	4	8.2	
1995	4.0	8	17.0	
4156	2.0	4	10.1	
6328	1.5	4	7.8	

	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)	\
6307	8.5	9.4	
5036	6.8	7.6	
1995	9.9	13.8	
4156	7.8	9.1	
6328	6.5	7.2	

	Fuel Consumption Comb (mpg)
6307	30
5036	37
1995	20
4156	31

```
6328          39
6307      216
5036      177
1995      321
4156      209
6328      168
Name: CO2 Emissions(g/km), dtype: int64
```

✓ Random Forest

The optimized model used was the random forest model. A random forest model is a supervised machine learning algorithm. It “combines the output of multiple decision trees to reach a single result.

It’s use and flexibility have also fueled its adoption,” (IBM, n.d.).

The table below summarizes the performance metrics of the random forest model and its accuracy.

Metric	Value
Random Forest Accuracy	0.5396073121191605
Random Forest Mean Squared Error (MSE)	41.3696682464455
Random Forest R-squared (R^2)	0.9884923876429129
Random Forest Root Mean Squared Error (RMSE)	6.4319257028082575

Although the accuracy of the random forest model is not that high, its performance metrics speaks for its excellent performance.

The MSE “measures the amount of error in statistical models. It assesses the average squared difference between the observed and predicted values. When a model has no error, the MSE equals zero,” (Statistics by Jim, n.d.). An MSE of 41.37 is considered low, indicating the predicted values are quite accurate in terms of deviation from the actual values.

An RMSE “is one of the two main performance indicators for a regression model.

It measures the average difference between values predicted by a model and the actual values. It provides an estimation of how well the model is able to predict the target value (accuracy),” (SAP, n.d.).

An RMSE of 6.43 is considered low, indicating the predicted values are quite accurate in terms of deviation from the actual values.

Lastly, we look at the r-squared value. The r-squared value of 0.99 is a very positive indicator for the model.

An r-squared value “shows how well the data fit the regression model (the goodness of fit,” (Taylor, n.d.).

An r-squared of 0.99 suggests a very strong correlation between predicted and actual values.

R-squared ranges from 0 to 1, where when it is closer to one it suggests that there is more variance in the target variable captured by the model. An r-squared value of 0.99 is very high and one can conclude that the model fits the data extremely well.

Essentially, with the exception of the accuracy, the performance metrics indicate a well-performing model.

To compare with the results of the initial model, the table below shows what the initial model's performance metrics were.

```
# Importing RandomForestClassifier and metrics modules from scikit-learn
from sklearn.ensemble import RandomForestClassifier
from sklearn import metrics
```

```
# Create a Random Forest Classifier
rf_model = RandomForestClassifier(n_estimators=100, random_state=0)
```

```
# Train the model on the training data
rf_model.fit(X_train, y_train)
```

```
# Predict on the test set
y_pred_rf = rf_model.predict(X_test)
```

```
# Calculate accuracy (or any other desired evaluation metric)
accuracy_rf = metrics.accuracy_score(y_test, y_pred_rf)
print("Random Forest Accuracy:", accuracy_rf)
```

```
➦ Random Forest Accuracy: 0.5396073121191605
```

```
# Importing necessary libraries
import pandas as pd
from sklearn.ensemble import RandomForestClassifier
from sklearn import metrics
import numpy as np

# Assuming you have already trained the model with the given code
# rf_model = RandomForestClassifier(n_estimators=100, random_state=0)
# rf_model.fit(X_train, y_train)

# Create a new DataFrame 'new_data' with the same features as 'X'
new_data = pd.DataFrame([[0] * len(X.columns)], columns=X.columns)

# Assuming 'new_data' has the same features as 'X', you can make predictions as follows:
y_pred_new = rf_model.predict(new_data)

# The 'y_pred_new' variable now contains the predicted classes for the new data.

# If you are using the RandomForestClassifier for regression, you can predict continuous
y_pred_new_regression = rf_model.predict(new_data)

# The 'y_pred_new_regression' variable now contains the predicted values for the new data

# Evaluating the model on test data
# Assuming you already have y_test (the true target values) and y_pred_rf (predicted values)

# Calculate Mean Squared Error (MSE) on the test set
mse_rf = metrics.mean_squared_error(y_test, y_pred_rf)

# Calculate R-squared ( $R^2$ ) on the test set
r_squared_rf = metrics.r2_score(y_test, y_pred_rf)

# Calculate Root Mean Squared Error (RMSE) on the test set
rmse_rf = np.sqrt(mse_rf)

print("Random Forest Mean Squared Error (MSE):", mse_rf)
print("Random Forest R-squared ( $R^2$ ):", r_squared_rf)
print("Random Forest Root Mean Squared Error (RMSE):", rmse_rf)
```



```
Random Forest Mean Squared Error (MSE): 41.3696682464455
Random Forest R-squared ( $R^2$ ): 0.9884923876429129
Random Forest Root Mean Squared Error (RMSE): 6.4319257028082575
```

```
# Importing necessary libraries
import pandas as pd
from sklearn.ensemble import RandomForestClassifier
from sklearn import metrics

# Assuming you have already trained the model with the given code
# rf_model = RandomForestClassifier(n_estimators=100, random_state=0)
# rf_model.fit(X_train, y_train)

# Create a new DataFrame 'new_data' with the same features as 'X'
new_data = pd.DataFrame(columns=X.columns)

# Add a single dummy row to 'new_data' using pandas.concat()
new_row = pd.DataFrame([[0] * len(X.columns)], columns=X.columns)
new_data = pd.concat([new_data, new_row], ignore_index=True)

# Assuming 'new_data' has the same features as 'X', you can make predictions as follows:
y_pred_new = rf_model.predict(new_data)


# The 'y_pred_new' variable now contains the predicted classes for the new data.

# If you are using the RandomForestClassifier for regression, you can predict continuous
y_pred_new_regression = rf_model.predict(new_data)

# The 'y_pred_new_regression' variable now contains the predicted values for the new data

# Show the predicted output for classification (predicted classes)
print("Predicted Classes for New Data:")
print(y_pred_new)

# Show the predicted output for regression (predicted values)
print("Predicted Values for New Data:")
print(y_pred_new_regression)
```

 Predicted Classes for New Data:
[96]
Predicted Values for New Data:
[96]

```
# Importing necessary libraries
import pandas as pd
from sklearn.ensemble import RandomForestClassifier
from sklearn import metrics

# Assuming you have already trained the model with the given code
# rf_model = RandomForestClassifier(n_estimators=100, random_state=0)
# rf_model.fit(X_train, y_train)

# Load your new data with feature values for over 1,400 data points.
# Ensure that 'new_data' has the same features as 'X_train' (excluding the target column)
# Replace this with the actual data containing feature values for each data point you want
new_data = pd.read_csv("/content/drive/MyDrive/final_dataset.csv")

# Drop the target column 'CO2 Emissions(g/km)' from 'new_data' if it's present
if 'CO2 Emissions(g/km)' in new_data.columns:
    new_data.drop(columns=['CO2 Emissions(g/km)'], inplace=True)

# Drop the columns that were dropped from the training data
columns_to_drop = ['uuid', 'Make', 'Vehicle Class', 'Transmission', 'Fuel Type', 'Model']
new_data.drop(columns=columns_to_drop, inplace=True)

# Assuming 'new_data' has the same features (except the dropped ones and target column) as
y_pred_new = rf_model.predict(new_data)

# The 'y_pred_new' variable now contains the predicted classes for each data point in 'new_data'

# Show the predicted output for classification (predicted classes)
print("Predicted Classes for New Data:")
for prediction in y_pred_new:
    print(prediction)
```



248
193
191
196
137
305
258
186
271
307
313
296
196
222
241
312
228
303
432
259
210
356
197
235
209
325
198
339
283
196
210
180
293
301

✓ Prediction of CO2 Emissions

```
# Importing necessary libraries
import pandas as pd
from sklearn.ensemble import RandomForestRegressor
from sklearn import metrics

new_data = pd.read_csv("/content/drive/MyDrive/final_dataset.csv")

# Drop the target column 'CO2 Emissions(g/km)' from 'new_data' if it's present
if 'CO2 Emissions(g/km)' in new_data.columns:
    new_data.drop(columns=['CO2 Emissions(g/km)'], inplace=True)

# Drop the columns that were dropped from the training data
columns_to_drop = ['uuid', 'Make', 'Vehicle Class', 'Transmission', 'Fuel Type', 'Model']
new_data.drop(columns=columns_to_drop, inplace=True)

# Assuming 'new_data' has the same features (except the dropped ones and target column) a
# and you have standardized the data, you can make predictions as follows:
y_pred_new_regression = rf_model.predict(new_data)

# Show the predicted output for regression (predicted CO2 values) and the true CO2 values
print("Predicted and True CO2 Values for New Data:")
for i, y_pred in enumerate(y_pred_new_regression):
    if i < len(y_test):
        print(f"Sample {i + 1}: Predicted CO2 Emissions = {y_pred}, True CO2 Emissions =
```



```
Sample 1451: Predicted CO2 Emissions = 271, True CO2 Emissions = 303
Sample 1452: Predicted CO2 Emissions = 230, True CO2 Emissions = 267
Sample 1453: Predicted CO2 Emissions = 219, True CO2 Emissions = 189
Sample 1454: Predicted CO2 Emissions = 258, True CO2 Emissions = 239
Sample 1455: Predicted CO2 Emissions = 240, True CO2 Emissions = 409
Sample 1456: Predicted CO2 Emissions = 268, True CO2 Emissions = 330
Sample 1457: Predicted CO2 Emissions = 194, True CO2 Emissions = 200
```

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestRegressor

new_data = pd.read_csv("/content/drive/MyDrive/final_dataset.csv")

# Drop the target column 'CO2 Emissions(g/km)' from 'new_data' if it's present
if 'CO2 Emissions(g/km)' in new_data.columns:
    new_data.drop(columns=['CO2 Emissions(g/km)'], inplace=True)

# Drop the columns that were dropped from the training data
columns_to_drop = ['uuid', 'Make', 'Vehicle Class', 'Transmission', 'Fuel Type', 'Model']
new_data.drop(columns=columns_to_drop, inplace=True)

# Assuming 'new_data' has the same features (except the dropped ones and target column) as
# and you have standardized the data, you can make predictions as follows:
y_pred_new_regression = rf_model.predict(new_data)

# Create a DataFrame with the predicted CO2 values for easy plotting
results_df = pd.DataFrame({'Predicted CO2 Emissions': y_pred_new_regression})

# Reset the index of y_test to ensure proper alignment
y_test.reset_index(drop=True, inplace=True)

# Plotting the 'Predicted CO2' against the data index (as we don't have true CO2 values for
plt.scatter(range(len(results_df)), results_df['Predicted CO2 Emissions'], c='b', label='Pi

# Plotting the true CO2 values for the original data (assuming you still have y_test avail
plt.scatter(range(len(y_test)), y_test, c='r', label='True CO2 Emissions')
```